
Unsupervised object learning from dense equivariant image labelling

James Thewlis

Hakan Bilen

Andrea Vedaldi

Visual Geometry Group, University of Oxford
{jdt,hbilen,vedaldi}@robots.ox.ac.uk

Abstract

One of the key challenges of visual perception is to extract abstract models of 3D objects and object categories from visual measurements, which are affected by complex nuisance factors such as viewpoint, occlusion, motion, and deformations. Starting from the recent idea of viewpoint factorization, we propose a new approach that, given a large number of images of an object and no other supervision, can extract a dense object-centric coordinate frame. This coordinate frame is invariant to deformations of the images and comes with a dense equivariant labelling neural network that can map image pixels to their corresponding object coordinates. We demonstrate the applicability of this method to simple articulated objects and deformable objects such as human faces, learning embeddings from random synthetic transformations or optical flow correspondences, all without any manual supervision.

1 Introduction

Humans can easily construct mental models of complex 3D objects and object categories from visual observations. This is remarkable because the dependency of an object’s appearance to its structure is tangled in a complex manner with extrinsic nuisance factors such as viewpoint, illumination, and articulation. Therefore, learning the intrinsic structure of an object from images requires removing these unwanted factors of variation from the data.

The recent work of [30] has proposed an unsupervised approach to do so, based on the concept of *viewpoint factorization*. The idea is to learn a deep Convolutional Neural Network (CNN) that can, given an image of the object, detect a discrete set of object landmarks. Differently from traditional approaches to landmark detection, however, landmarks are not defined nor supervised manually. Instead, the detectors are learned using the only requirement that the detected points must be equivariant (consistent) with deformations of the input images. The authors of [30] show that this constraint is sufficient to learn landmarks that are “intrinsic” to the objects and hence capture their structure; remarkably, due to the generalization ability of CNNs, the landmark points are detected consistently not only across deformations of a given object instance, which are observed during training, but also across different instances. This behaviour emerges automatically from training on thousands of single-instance correspondences.

In this paper, we take this idea further, moving beyond a sparse set of landmarks to a dense model of the object structure (section 3). Our method relates each point on an object to a point in a low dimensional vector space in a way that is consistent across variation in motion and in instance identity. This gives rise to an object-centric coordinate system, which allows points on the surface of an object to be indexed semantically (figure 1). As an illustrative example, take the object category of a face and the vector space \mathbb{R}^3 . Our goal is to semantically map out the object such that any point on a face, such as the left eye, lives at a canonical position in this “label space”. We train a CNN to learn the function that projects any face image into this space, essentially “coloring” each pixel with its corresponding label. As a result of our learning formulation, the label space has the property of being

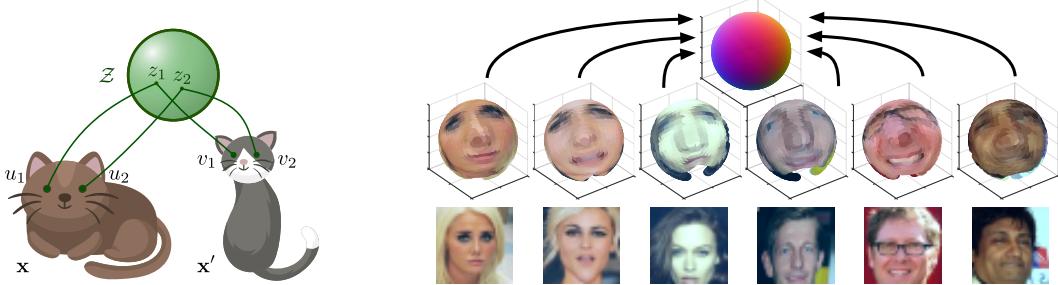


Figure 1: **Dense equivariant image labelling.** *Left:* Given a large images x of an object or object category and no other supervision, our goal is to find a common latent space \mathcal{Z} , homeomorphic to a sphere, which attaches a semantically-consistent coordinate frame to the object points. This is done by learning a dense labelling function that maps image pixels to their corresponding coordinate in the \mathcal{Z} space. This mapping function is equivariant (compatible) with image warps or object instance variations. *Right:* An equivariant dense mapping learned in an unsupervised manner from a large dataset of faces.

locally smooth: points nearby in the image are nearby in the label space. In an ideal case, we could imagine the surface of an object to be mapped to a sphere.

In order to achieve these results, we contribute several technical innovations (section 3.2). First, we show that, in order to learn a non-trivial object coordinate frame, the concept of equivariance must be complemented with the one of *distinctiveness* of the embedding. Then, we propose a CNN implementation of this concept that can explicitly express uncertainty in the labelling of the object points. The formulation is used in combination with a probabilistic loss, which is augmented with a robust geometric distance to encourage better alignment of the object features.

We show that this framework can be used to learn meaningful object coordinate frames in a purely unsupervised manner, by analyzing thousands of deformations of visual objects. While [30] proposed to use synthetic image warps for training, here we also consider simple articulated objects having frames related by optical flow (section 4).

We conclude the paper with a summary of our finding (section 5).

2 Related Work

Learning the structure of visual objects. Modeling the structure of visual objects is a widely-studied (*e.g.* [4, 5, 8]) computer vision problem with important applications such as facial landmark detection and human body pose estimation. Most of this work is supervised and aimed at learning detectors of the objects and their parts, often using deep learning. A few approaches such as spatial transformer networks [16] can learn geometric transformations without explicit geometric supervision, but do not build explicit geometric models of visual objects.

More related to our work, WarpNet [17] learns a neural network that predicts a Thin Plate Spline transformations between pairs of image of an object. Deep Deformation Network [31] improves WarpNet by using Point Transformer Network to refine the computed landmarks, however, it requires manual supervision. However, none of these works look at the problem of learning an invariant geometric embedding for the object.

Our work build on the idea of *viewpoint factorization* (section 3.1), recently introduced in [30]. However, we extend [30] in several significant ways. First, we construct a *dense* rather than discrete embedding, where all pixels of an object are mapped to an invariant object-centric coordinate instead of just a small set of selected landmarks. Second, we show that the equivariance constraint proposed in [30] is not quite enough to learn such an embedding; it must be complemented with the concept of a *distinctive* embedding (section 3.1). Third, we introduce a new neural network architecture and corresponding training objective that allow such an embedding to be learned in practice (section 3.2).

Optical/semantic flow. A common technique to find correspondences between temporally related video frames is computing optical flow [14]. The state-of-the-art methods [10, 29, 15] typically employ convolutional neural networks to learn pairwise dense correspondences between the same

object instances at subsequent frames. The SIFT Flow method [20] extends the between-instance correspondences to cross-instance mappings by matching SIFT features [22] between semantically similar object instances. Learned-Miller [19] extends the pairwise correspondences to multiple images by posing a problem of alignment among the images of a set. Collection Flow [18] and Mobahi *et al.* [24] project objects onto a low-rank space that allow for joint alignment. FlowWeb [36], and Zhou *et al.* [35] construct fully connected graphs to maximise cycle consistency between each image pair and synthetic data as an intermediary by training a CNN. As a matter of fact, our work can build on optical flow as a supervisory signal (section 4), extracting from it the invariant object factor.

Unsupervised learning. Classical unsupervised learning methods such as autoencoders [3, 2, 13] and denoising autoencoders aim to learn useful feature representations from an input by simply reconstructing it after a bottleneck. In addition to the autoencoders, the generative adversarial networks [12] target producing samples of realistic images by training generative models. These models when trained joint with image encoders are also shown to learn good feature representations [7]. More recently several studies have emerged that train neural networks by learning auxiliary or pseudo tasks. These methods exploit typically some existing information in input as “self-supervision” without any manual labeling by removing or perturbing some information from an input and requiring a network to reconstruct it. For instance, Doersch *et al.* [6], and Noroozi and Favaro [25] train a network to predict the relative locations of shuffled image patches. Other self-supervised tasks include colorizing images [32], inpainting [27], ranking frames of a video in temporally correct order [23, 9]. More related to our approach, Agrawal *et al.* [1] use egomotion as supervisory signal to learn feature representations in a Siamese network by predicting camera transformations from image pairs, [26] learn to group pixels that move together in a video. [34, 11] use a warping-based loss to learn depth from video.

3 Method

This section discusses our method in detail, first introducing the general idea of dense equivariant labelling (section 3.1), and then presenting a concrete implementation of the latter using a novel deep CNN architecture (section 3.2).

3.1 Dense equivariant labelling

Consider a 3D object $S \subset \mathbb{R}^3$ or a class of such objects S that are topologically isomorphic to a sphere $\mathcal{Z} \subset \mathbb{R}^3$ (i.e. the objects are simple closed surfaces without holes). We can construct a homeomorphism $p = \pi_S(q)$ mapping points of the sphere $q \in \mathcal{Z}$ to points $p \in S$ of the objects. Furthermore, if the objects belong to the same semantic category (e.g. faces), we can assume that these isomorphisms are *semantically consistent*, in the sense that $\pi_{S'} \circ \pi_S^{-1} : S \rightarrow S'$ maps points of object S to semantically-analogous points in object S' (e.g. for human faces the right eye in one face should be mapped to the right eye in another [30]).

While this construction is abstract, it shows that we can endow the object (or object category) with a spherical reference system \mathcal{Z} . The authors of [30] build on this construction to define a discrete system of object landmarks by considering a finite number of points $z_k \in \mathcal{Z}$. Here, we take the geometric embedding idea more literally and propose to explicitly learn a dense mapping from images of the object to the object-centric coordinate space \mathcal{Z} . Formally, we wish to learn a *labelling function* $\Phi : (\mathbf{x}, u) \mapsto z$ that takes a RGB image $\mathbf{x} : \Lambda \rightarrow \mathbb{R}^3$, $\Lambda \subset \mathbb{R}^3$ and a pixel $u \in \Lambda$ to the object point $z \in \mathcal{Z}$ which is imaged at u (figure 1).

Similarly to [30], this mapping must be compatible or equivariant with image deformations. Namely, let $g : \Lambda \rightarrow \Lambda$ be a deformation of the image domain, either synthetic or due to a viewpoint change or other motion. Furthermore, let $g\mathbf{x} = \mathbf{x} \circ g^{-1}$ be the action of g on the image (obtained by inverse warp). Barring occlusions and boundary conditions, we must have the following *equivariance constraint*:

$$\forall \mathbf{x}, u : \quad \Phi(\mathbf{x}, u) = \Phi(g\mathbf{x}, gu). \quad (1)$$

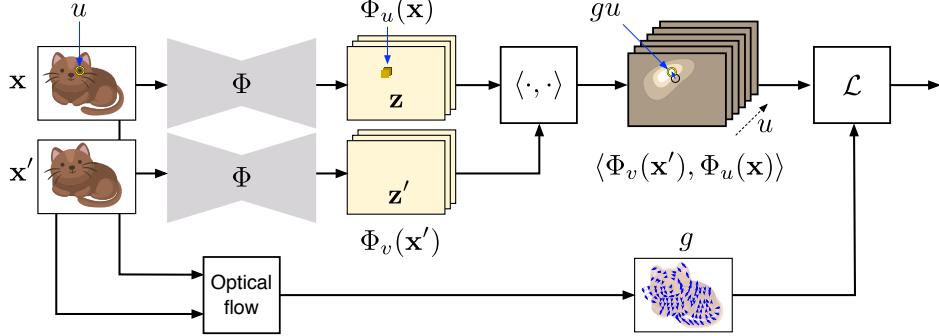


Figure 2: **Unsupervised dense correspondence network.** From left to right: The network Φ extracts label maps $\Phi_u(\mathbf{x})$ and $\Phi_v(\mathbf{x}')$ from the image pair \mathbf{x} and \mathbf{x}' . An optical flow module (or ground truth for synthetic transformation) computes the warp (correspondence field) g such that $\mathbf{x}' = g\mathbf{x}$. Then the label of each point u in the first image is correlated to each point v in the second, obtaining a number of score maps. The loss evaluates how well the score maps predict the warp g .

This constraint simply states that pixel u in image \mathbf{x} must receive the same label as pixel gu in image $g\mathbf{x}$. Naively, we may translate this constraint into the loss function

$$\mathcal{L}(\Phi|\alpha) = \frac{1}{|\Lambda|} \int_{\Lambda} \|\Phi(\mathbf{x}, u) - \Phi(g\mathbf{x}, gu)\|^2 du.$$

However, minimizing this loss has the significant drawback that a global optimum is obtained by simply setting $\Phi(\mathbf{x}, u) = \text{const}$. The reason for this issue is that (1) is not quite enough to learn a useful object representation. In order to do so, we must require the labels not only to be equivariant, but also *distinctive*, in the sense that

$$\Phi(\mathbf{x}, u) = \Phi(g\mathbf{x}, v) \Leftrightarrow v = gu.$$

We can encode this requirement as a loss in different ways. For example, by using the fact that points $\Phi(\mathbf{x}, u)$ are on the unit sphere, we can use the loss:

$$\mathcal{L}'(\Phi|\mathbf{x}, g) = \frac{1}{|\Lambda|} \int_{\Lambda} \|gu - \text{argmax}_v \langle \Phi(\mathbf{x}, u), \Phi(g\mathbf{x}, v) \rangle\|^2 du. \quad (2)$$

By doing so, the labels $\Phi(\mathbf{x}, u)$ must be able to discriminate between different object points, so that a constant labelling would receive a high penalty.

Relationship with learning invariant visual descriptors. As an alternative to loss (2), we could have used a pairwise loss¹ to encourage the similarity $\langle \Phi(\mathbf{x}, u), \Phi(\mathbf{x}', gu) \rangle$ of the labels assigned to corresponding pixels u and gu to be larger than the similarity $\langle \Phi(\mathbf{x}, u), \Phi(\mathbf{x}', v) \rangle$ of the labels assigned to pixels u and v that do *not* correspond. Formally, this would result in a pairwise loss similar to the ones often used to learn invariant visual descriptors for image matching. The reason why our method learns an object representation instead of a generic visual descriptor is that the *dimensionality* of the label space \mathcal{Z} is just enough to represent a point on a surface. If we replace \mathcal{Z} with a larger space such as \mathbb{R}^d , $d \gg 2$, we can expect $\Phi(\mathbf{x}, u)$ to learn to extract generic visual descriptors such as SIFT instead. This establishes an interesting relationship between visual descriptors and object-specific coordinate vectors and suggests that it is possible to transition between the two by controlling their dimensionality.

3.2 Concrete learning formulation

In this section we introduce a concrete implementation of our method (figure 2). For the mapping Φ , we use a CNN that receives as input an image tensor $\mathbf{x} \in \mathbb{R}^{H \times W \times C}$ and produces as output a

¹Formally, this is achieved by the loss

$$\mathcal{L}''(\Phi|\mathbf{x}, g) = \frac{1}{|\Lambda|} \int_{\Lambda} \max \left\{ 0, \max_v \Delta(u, v) + \langle \Phi(\mathbf{x}, u), \Phi(g\mathbf{x}, v) \rangle - \langle \Phi(\mathbf{x}, u), \Phi(g\mathbf{x}, gu) \rangle \right\} du,$$

where $\Delta(u, v) \geq 0$ is an error-dependent margin.

label tensor $\mathbf{z} \in \mathbb{R}^{H \times W \times L}$. We use the notation $\Phi_u(\mathbf{x})$ to indicate the L -dimensional label vector extracted at pixel u from the label image computed by the network.

The dimension of the label vectors is set to $L = 3$ (instead of $L = 2$) in order to allow the network to express uncertainty about the label assigned to a pixel. The network can do so by modulating the norm of $\Phi_u(\mathbf{x})$. In fact, correspondences are expressed probabilistically by computing the inner product of label vectors followed by the softmax operator. Formally, the probability that pixel v in image \mathbf{x}' corresponds to pixel u in image \mathbf{x} is expressed as:

$$p(v|u; \mathbf{x}, \mathbf{x}', \Phi) = \frac{e^{\langle \Phi_u(\mathbf{x}), \Phi_v(\mathbf{x}') \rangle}}{\sum_z e^{\langle \Phi_u(\mathbf{x}), \Phi_z(\mathbf{x}') \rangle}}. \quad (3)$$

In this manner, a shorter vector Φ_u results in a more diffuse probability distribution.

Next, we wish to define a loss function for learning Φ from data. To this end, we consider triplet $\alpha = (\mathbf{x}, \mathbf{x}', g)$, where $\mathbf{x}' = g\mathbf{x}$ is an image that corresponds to \mathbf{x} up to transformation g (the nature of the data is discussed below). We then assess the performance of the network Φ on the tripled α using two losses. The first loss is the *negative log-likelihood* of the ground-truth correspondences:

$$\mathcal{L}_{\log}(\Phi|\mathbf{x}, \mathbf{x}', g) = -\frac{1}{HW} \sum_u \log p(gu|u; \mathbf{x}, \mathbf{x}', \Phi). \quad (4)$$

This loss has the advantage that it explicitly learns (3) as the probability of a match. However, it is not sensitive to the *size* of a correspondence error $v - gu$. In order to address this issue, we also consider the loss

$$\mathcal{L}_{\text{dist}}(\Phi|\mathbf{x}, \mathbf{x}', g) = \frac{1}{HW} \sum_u \sum_v \|v - gu\|_2^\gamma p(v|u; \mathbf{x}, \mathbf{x}', \Phi). \quad (5)$$

Here $\gamma > 0$ is an exponent used to control the robustness of the distance measure, which we set to $\gamma = 0.5, 1$.

The structure of our network is fairly standard, based on that used in [33] and adapted in [30], having convolutional filter sizes $[20, 48, 64, 80, 256, L]$ and a single max pooling of stride 2 after the first convolutional layer. Consequently, our \mathbf{z} is in practice downsampled by 2.

3.3 Learning from synthetic and real deformations

Losses (4) and (5) learn from triplets $\alpha = (\mathbf{x}, \mathbf{x}', g)$. Here \mathbf{x}' can be either generated synthetically by applying a random transformation g to a natural image \mathbf{x} [30, 17], or it can be obtained by observing image pairs $(\mathbf{x}, \mathbf{x}')$ containing real object deformations arising from a viewpoint change or an object motion or deformation.

The use of synthetic transformations was considered in [30], who showed it to be sufficient to learn meaningful landmarks for a number of real-world objects such as human and cat faces. Here, in addition to using synthetic deformations, we also consider using real-world image pairs \mathbf{x} and \mathbf{x}' . In principle, the learning formulation can be modified so that knowledge of g is not required; instead, images and their warps can be compared and aligned directly based on the brightness constancy principle. However, for simplicity and robustness, we prefer to pre-process \mathbf{x} and \mathbf{x}' by using an off-the-shelf optical flow algorithm; this produces a noisy version of g which, when used as ground-truth in our learning formulation, may still be good enough to learn meaningful object embeddings.

4 Experiments

This section assesses our unsupervised method for dense object labelling on two representative tasks: a toy problem (section 4.1) and human faces (section 4.2).

4.1 Toy example

In order to illustrate our method we consider a toy problem consisting of a simple articulated object, namely an animated robotic arm (figure 3) created using a 2D physics engine [28]. We do so for two reasons: to show that the approach is capable of labelling correctly deformable/articulated objects and to show that the spherical model \mathcal{Z} is applicable also to thin objects, that have mainly a 1D structure.

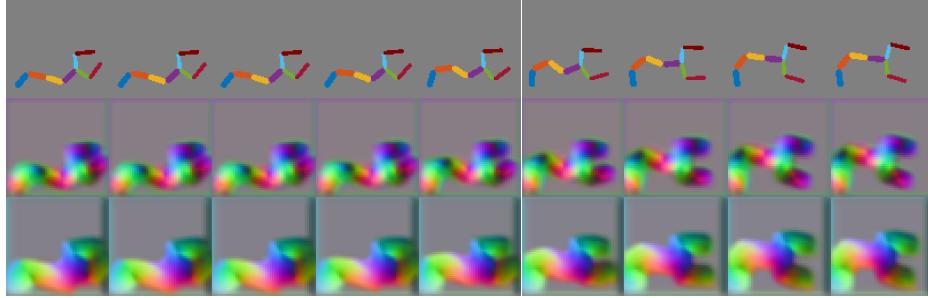


Figure 3: **Roboarm equivariant labelling.** Top: Original video frames of a simple articulated object. Middle and bottom row: color-coded spherical embedding of the object obtained using the probabilistic and distance-based losses, respectively.

Dataset details. The arm is anchored to the bottom left corner and is made up of colored capsules connected with joints having reasonable angle limits to prevent unrealistic contortion and self-occlusion. Motion is achieved by varying the gravity vector, sampling each element from a Gaussian with standard deviation 15 m s^{-2} every 100 iterations. Frames \mathbf{x} of size 90×90 pixels and the corresponding flow fields $g : \mathbf{x} \mapsto \mathbf{x}'$ are saved every 20 iterations. We also save the positions of the capsule centers. The final dataset has 23999 frames.

Learning. Using the correspondences $\alpha = (\mathbf{x}, \mathbf{x}', g)$ provided by the flow fields, we use our method to learn an object centric coordinate frame \mathcal{Z} and its corresponding labelling function $\Phi_u(\mathbf{x})$. We test learning Φ using the probabilistic loss (4) and distance-based loss (5). In the loss we ignore areas with zero flow, which automatically removes the background.

Results. Figure 3 provides some qualitative results, showing by means of colormaps the labels $\Phi_u(\mathbf{x})$ associated to different pixels of each input image. It is easy to see that the method attaches consistent labels to the different arm elements. The distance-based loss produces a more uniform embedding, as it may be expected. The embeddings are further visualized in Figure 4 by projecting a number of video frames back to the learned coordinate spaces \mathcal{Z} . It can be noted that the space is invariant, in the sense that the resulting figure is approximately the same despite the fact that the object deforms significantly in image space. This is true for both embeddings, but the distance-based ones are geometrically more consistent.

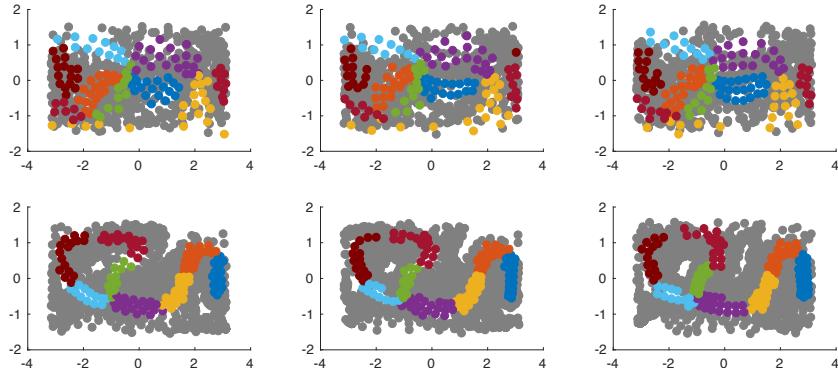


Figure 4: **Invariance of the object-centric coordinate space for Roboarm.** The plot projects frames 3,6,9 of figure 3 on the object-centric coordinate space \mathcal{Z} , using the embedding functions learned by means of the probabilistic (top) and distance (bottom) based losses. The sphere is then unfolded, plotting latitude and longitude (in radians) along the vertical and horizontal axes.

Predicting capsule centers. We evaluate quantitatively the ability of our *object frames* to be used for localization of the capsule centers. If our theory is correct and a coordinate system intrinsic to the object has been learned, then we should expect there to be a specific 3-vector in \mathcal{Z} corresponding to each center, and our job is to find these vectors. Various strategies could be used, such as averaging the object-centric coordinates given to the centers over the training set, but we choose to incorporate the problem into the learning framework. This is done using the negative log-likelihood in much

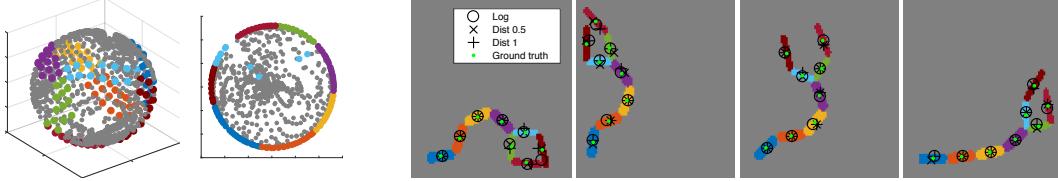


Figure 5: Left: Embedding spaces of different dimension. Spherical embedding (from the 3D embedding function $\Phi_u(\mathbf{x}) \in \mathbb{R}^3$) learned using the distance loss compared to a version with one dimension less. **Right: Capsule center prediction for different losses.**

the same way as (4), limiting our vectors u to the centers. This is done as an auxiliary layer with no backpropagation to the rest of the network, so that the embedding remains unsupervised. The error reported is the Euclidean distance as a percentage of the image width.

Results are given for the different choices of loss function used for unsupervised training in Table 1 and visualized in Figure 5, showing that the object centers can be located to a high degree of accuracy. The negative log likelihood performs best while the two losses incorporating distance perform similarly.

We also perform experiments varying the dimensionality L of the label space \mathcal{Z} (Table 2). Perhaps most interestingly, given the almost one-dimensional nature of the arm, is the case of $L = 2$, which would correspond to a circular space (since the length of vectors is used to code for uncertainty). As can be seen in the right of Figure 5, the segments are represented almost perfectly on the boundary of a circle, with the exception of the bifurcation which it is unable to accurately represent. This is manifested by the light blue segment trying, and failing, to be in two places at once.

Unsupervised Loss	Error
\mathcal{L}_{log}	0.97 %
$\mathcal{L}_{\text{dist}}, \gamma = 1$	1.13 %
$\mathcal{L}_{\text{dist}}, \gamma = 0.5$	1.14 %

Table 1: Predicting capsule centers.
Error as percent of image width.

Descriptor Dimension	Error
2	1.29 %
3	1.14 %
5	1.16 %
20	1.28 %

Table 2: Descriptor dimension
($\mathcal{L}_{\text{dist}}, \gamma = 0.5$). $L > 3$ shows no improvement,
suggesting $L=3$ is the natural manifold of the arm.

4.2 Faces

After testing the properties of our method on a simple toy problem, we move to a much harder task and apply our method to generate an object-centric reference frame \mathcal{Z} for the *category* of human faces. In order to generate an image pair and corresponding flow field for training we warp each face synthetically using Thin Plate Spline warps in a manner similar to [30]. We train our models on the official training split of the extensive CelebA [21] dataset which contains 200k celebrity faces with annotations of 5 keypoints, the eyes, nose and mouth corners. Please note that We do not use any of these annotations to train our model.

Qualitative assessment. We find that the negative log-likelihood loss, while performing best for the simple example of the arm, performs poorly to model the face due to the variance in appearance and shape across different instances and more complicated deformable geometry. Specifically, this model fails to disambiguate the left and right eye, as shown in Figure 7. The distance-based loss (5) produces more coherent embedding, as can be noted in Figure 6.

Regressing semantic keypoints. We would like to quantify the accuracy of our model in terms of ability to consistently locate manually annotated points, specifically the eyes, nose, and mouth corners given in the CelebA dataset. We use the standard test split for evaluation of the MAFL dataset [33], containing 1000 images. We also use the MAFL training subset of 19000 images for learning to predict the ground truth keypoints, which gives a quantitative measure of the consistency of our *object frame* for detecting facial features. These are reported as Euclidean error normalized as a percentage of inter-ocular distance.

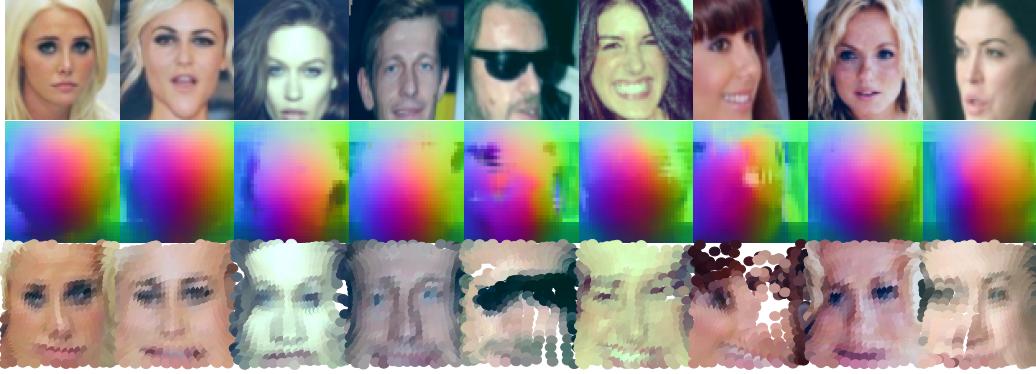


Figure 6: **Faces.** Top: Input images, Middle: Dimensions mapped to RGB, Bottom: Image pixels mapped to R3 and flattened as spherical coordinates



Figure 7: **Left:** Semantic keypoints regressed using our unsupervised face reference system. **Right:** failure of the log loss. Three image of ambiguously labelled eyes and poorly regressed keypoints. Ground truth in light blue.

In order to map the object frame to the semantic keypoints, as in the case of the robot arm centers, we learn the vectors $z_k \in \mathcal{Z}$ corresponding to the position of each point in our canonical reference space and then, for any given image, find the most similar z and its corresponding pixel location u . We report the localization performance of this model in Table 3. We empirically validate that the negative log-likelihood is not ideal for this task and obtain higher performance for the robust distance with power 0.5. However, we observe that the performance of our model is limited by the low spatial resolution due to subsampling along the network which can be improved by using a set of z_k .

The method of [30] learns to regress P ground truth coordinates based on $M > P$ unsupervised landmarks. By regressing from multiple points it is not limited to integer pixel coordinates. While we are not predicting landmarks as network output, we can emulate this method by allowing multiple points in our object coordinate space to be predictive for a single ground truth landmark. We learn one regressor per ground truth point, each formulated as a linear combination of $M = 50$ intermediate points. Essentially, this allows the network to say which points in \mathcal{Z} are most useful for predicting each ground truth point.

As shown in Table 4, we perform comparably to other unsupervised landmark prediction methods, despite the fact that we learn a full dense embedding of the object rather than a sparse set of points.

Unsupervised Loss	Error
\mathcal{L}_{log}	59.25 %
$\mathcal{L}_{\text{dist}}, \gamma = 1$	14.32 %
$\mathcal{L}_{\text{dist}}, \gamma = 0.5$	12.98 %

Table 3: Nearest neighbour keypoint prediction on CelebA

Unsupervised loss	Error
$\mathcal{L}_{\text{dist}}, \gamma = 1$	7.94 %
$\mathcal{L}_{\text{dist}}, \gamma = 0.5$	7.18 %
[30]	6.67 %

Table 4: Comparison on MAFL

5 Conclusions

Building on the idea of viewpoint factorization, we have introduce a new method that can endow an object or object category with an invariant dense geometric embedding automatically, by simply observing a large dataset of unlabelled images. Our learning framework combines in a novel way the concept of equivariance with the one of distinctiveness. We have also proposed a concrete implementation using novel losses to learn a deep dense image labeller. We have shown empirically that the method can learn consistent geometric embedding for a simple articulated synthetic robotic arm as well as for real human faces. The resulting embeddings are invariant to deformations and, importantly, to *intra-category* variations.

References

- [1] Pulkit Agrawal, Joao Carreira, and Jitendra Malik. Learning to see by moving. In *Proc. ICCV*, 2015.
- [2] Yoshua Bengio. Learning deep architectures for AI. *Foundations and trends in Machine Learning*, 2009.
- [3] H Bourlard and Y Kamp. Auto-Association by Multilayer Perceptrons and Singular Value Decomposition. *Biological Cybernetics*, 1988.
- [4] T F Cootes, C J Taylor, D H Cooper, and J Graham. Active shape models: their training and application. *CVIU*, 1995.
- [5] Navneet Dalal and Bill Triggs. Histograms of Oriented Gradients for Human Detection. In *Proc. CVPR*, 2005.
- [6] Carl Doersch, Abhinav Gupta, and Alexei A Efros. Unsupervised Visual Representation Learning by Context Prediction. In *Proc. ICCV*, 2015.
- [7] Jeff Donahue, Philipp Krähenbühl, and Trevor Darrell. Adversarial feature learning. *Proc. ICLR*, 2017.
- [8] Pedro F. Felzenszwalb, Ross B. Girshick, David McAllester, and Deva Ramanan. Object Detection with Discriminatively Trained Part Based Models. *PAMI*, 2010.
- [9] Basura Fernando, Hakan Bilen, Efstratios Gavves, and Stephen Gould. Self-supervised video representation learning with odd-one-out networks. In *Proc. CVPR*, 2017.
- [10] Philipp Fischer, Alexey Dosovitskiy, Eddy Ilg, Philip Häusser, Caner Hazırbaş, Vladimir Golkov, Patrick van der Smagt, Daniel Cremers, and Thomas Brox. FlowNet: Learning Optical Flow with Convolutional Networks. In *Proc. ICCV*, 2015.
- [11] Ravi Garg, Gustavo Carneiro, and Ian Reid. Unsupervised cnn for single view depth estimation: Geometry to the rescue. In *Proc. ECCV*, pages 740–756. Springer, 2016.
- [12] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016. <http://www.deeplearningbook.org>.
- [13] G E Hinton and R R Salakhutdinov. Reducing the Dimensionality of Data with Neural Networks. *Science*, 2006.
- [14] Berthold K.P. Horn and Brian G. Schunck. Determining optical flow. *Artificial Intelligence*, 1981.
- [15] Eddy Ilg, Nikolaus Mayer, Tonmoy Saikia, Margret Keuper, Alexey Dosovitskiy, and Thomas Brox. FlowNet 2.0: Evolution of Optical Flow Estimation with Deep Networks. *arXiv preprint arXiv:1612.01925*, 2016.
- [16] Max Jaderberg, Karen Simonyan, Andrew Zisserman, and Koray Kavukcuoglu. Spatial Transformer Networks. In *Proc. NIPS*, 2015.
- [17] A. Kanazawa, D. W. Jacobs, and M. Chandraker. WarpNet: Weakly supervised matching for single-view reconstruction. In *Proc. CVPR*, 2016.
- [18] Ira Kemelmacher-Shlizerman and Steven M. Seitz. Collection flow. In *Proc. CVPR*, 2012.
- [19] Erik G Learned-Miller. Data driven image models through continuous joint alignment. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 28(2):236–250, 2006.
- [20] Ce Liu, Jenny Yuen, and Antonio Torralba. SIFT Flow: Dense correspondence across scenes and its applications. *PAMI*, 2011.
- [21] Ziwei Liu, Ping Luo, Xiaogang Wang, and Xiaoou Tang. Deep learning face attributes in the wild. In *Proc. ICCV*, 2015.
- [22] David G Lowe. Distinctive image features from scale-invariant keypoints. *International journal of computer vision*, 60(2):91–110, 2004.
- [23] Ishan Misra, C Lawrence Zitnick, and Martial Hebert. Shuffle and learn: unsupervised learning using temporal order verification. In *Proc. ECCV*, 2016.

- [24] Hossein Mobahi, Ce Liu, and William T. Freeman. A Compositional Model for Low-Dimensional Image Set Representation. *Proc. CVPR*, 2014.
- [25] Mehdi Noroozi and Paolo Favaro. Unsupervised learning of visual representations by solving jigsaw puzzles. In *Proc. ECCV*, 2016.
- [26] Deepak Pathak, Ross Girshick, Piotr Dollár, Trevor Darrell, and Bharath Hariharan. Learning features by watching objects move. In *CVPR*, 2017.
- [27] Deepak Pathak, Philipp Krahenbuhl, Jeff Donahue, Trevor Darrell, and Alexei A Efros. Context Encoders: Feature Learning by Inpainting. In *Proc. CVPR*, 2016.
- [28] Yuval Tassa. CapSim - the MATLAB physics engine. <https://mathworks.com/matlabcentral/fileexchange/29249-capsim-the-matlab-physics-engine>.
- [29] J Thewlis, S Zheng, P. H. S. Torr, and A Vedaldi. Fully-Trainable Deep Matching. In *Proc. BMVC*, 2016.
- [30] James Thewlis, Hakan Bilen, and Andrea Vedaldi. Unsupervised learning of object landmarks by factorized spatial embeddings. *arXiv preprint arXiv:1705.02193*, 2017.
- [31] Xiang Yu, Feng Zhou, and Manmohan Chandraker. Deep Deformation Network for Object Landmark Localization. In Bastian Leibe, Jiri Matas, Nicu Sebe, and Max Welling, editors, *Proc. ECCV*, Cham, 2016.
- [32] Richard Zhang, Phillip Isola, and Alexei A Efros. Colorful Image Colorization. In *Proc. ECCV*, 2016.
- [33] Zhanpeng Zhang, Ping Luo, Chen Change Loy, and Xiaoou Tang. Learning Deep Representation for Face Alignment with Auxiliary Attributes. *PAMI*, 2016.
- [34] Tinghui Zhou, Matthew Brown, Noah Snavely, and David G Lowe. Unsupervised learning of depth and ego-motion from video. In *Proc. CVPR*, 2017.
- [35] Tinghui Zhou, Philipp Krähenbühl, Mathieu Aubry, Qixing Huang, and Alexei A. Efros. Learning Dense Correspondences via 3D-guided Cycle Consistency. In *Proc. CVPR*, 2016.
- [36] Tinghui Zhou, Yong Jae Lee, Stella X. Yu, and Alexei A. Efros. FlowWeb: Joint image set alignment by weaving consistent, pixel-wise correspondences. In *Proc. CVPR*, 2015.